


Lab 5: 网络层数据平面观察实验

part1



1. 第一个发送的UDP包

1. 发送端IP地址是172.17.19.70
2. 上层协议是17, 代表UDP
3. header length 20. 有效载荷Payload = total - header = 56 - 20 = 36.
4. 否。检查UDP包的Fragment Offset字段, 全都是0

2. 观察连续的UDP包 (穿插其他包)

1. 不断变化的有 Identification, Time to Live 和 Header Checksum.
保持不变的有

Version, Header Length, Differentiated Services Field, Total Length, Flags, Fragment Offset, Protocol, Source Address, Destination Address, Source Port, Destination Port.

```

  ...log
Internet Protocol Version 4, Src: 172.17.19.70, Dst: 202.120.224.81
0100 .... = Version: 4
.... 0101 = Header Length: 20 bytes (5)
Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
  0000 00.. = Differentiated Services Codepoint: Default (0)
  .... ..00 = Explicit Congestion Notification: Not ECN-Capable Transport (0)
Total Length: 56
Identification: 0xf36f (62319)
000. .... = Flags: 0x0
...0 0000 0000 0000 = Fragment Offset: 0
Time to Live: 1
Protocol: UDP (17)
Header Checksum: 0x5c24 [validation disabled]
```

```
[Header checksum status: Unverified]
Source Address: 172.17.19.70
Destination Address: 202.120.224.81
[Stream index: 1]
```

...

以及有效载荷的内容也不变，这是因为**traceroute**每次发送的UDP包的有效载荷内容是一样的。

2. 为什么有些字段不断变化，为什么有些不变？

必须保持不变：

版本：占4 bit，通信双方使用的版本必须一致，对于IPv4字段的值是4；

首部长度：占4 bit，首部长度说明首部有多少32位字（4字节）由于IPv4首部可能包含数目不定的选项，这个字段也用来确定数据的偏移量；

区分服务：占6bit，只有在使用区分服务时，这个字段才起作用，在一般的情况下都不使用这个字段；

保持不变：

显式拥塞通告：允许在不丢弃报文的同时通知对方网络拥塞的发生

全长：占16位字段，定义了报文总长，包含首部和数据，单位为字节。这个字段的最小值是20（0字节数据），最大值是65535。

标识符：占16位，这个字段主要被用来唯一地标识一个报文的所有分片，因为分片不一定按序到达，所以在重组时需要知道分片所属的报文

分片偏移：这个13位字段指明了每个分片相对于原始报文开头的偏移量，以8字节作单位。

源地址：报文的发送端；

目的地址：报文的接收端；

选项：附加的首部字段可能跟在目的地址之后；

必须更改

标识符：占16位，主要被用来唯一地标识一个报文的所有分片；

存活时间：占8位，避免报文在互联网中永远存在。实现为跳数计数器，报文经过的每个路由器都将此字段减1，当此字段等于0时，报文不再向下一跳传送并被丢弃，最大值是255

这是 **traceroute** 的核心原理；

首部检验和：占16位，检验和字段只对首部查错，在每一跳，路由器都要重新计算出的首部检验和并与此字段进行比对，如果不一致，此报文将会被丢弃；

数据：所谓的保持不变指的是这次 **traceroute** 不会改变的，但是下一次 **traceroute** 可能就会改了。

3. 列出连续IP数据报中的标识序列

0xf36f, 0x4609, 0xa333

3. 观察收到的第一个TTL-exceeded replies

1. 标识字段与TTL字段分别是多少？

0x0006, 128

2. 收到的所有TTL-exceeded replies中，这两个字段是否不变？为什么？

ID 改变，TTL 不变。因为每次收到的TTL-exceeded replies都是不同的报文，所以ID会改变，而TTL-exceeded replies是由同一个路由器返回的，到本机的转发次数相同（都是0），所以TTL不变。

part2

我的 Ubuntu 版本是 24.04.1 LTS

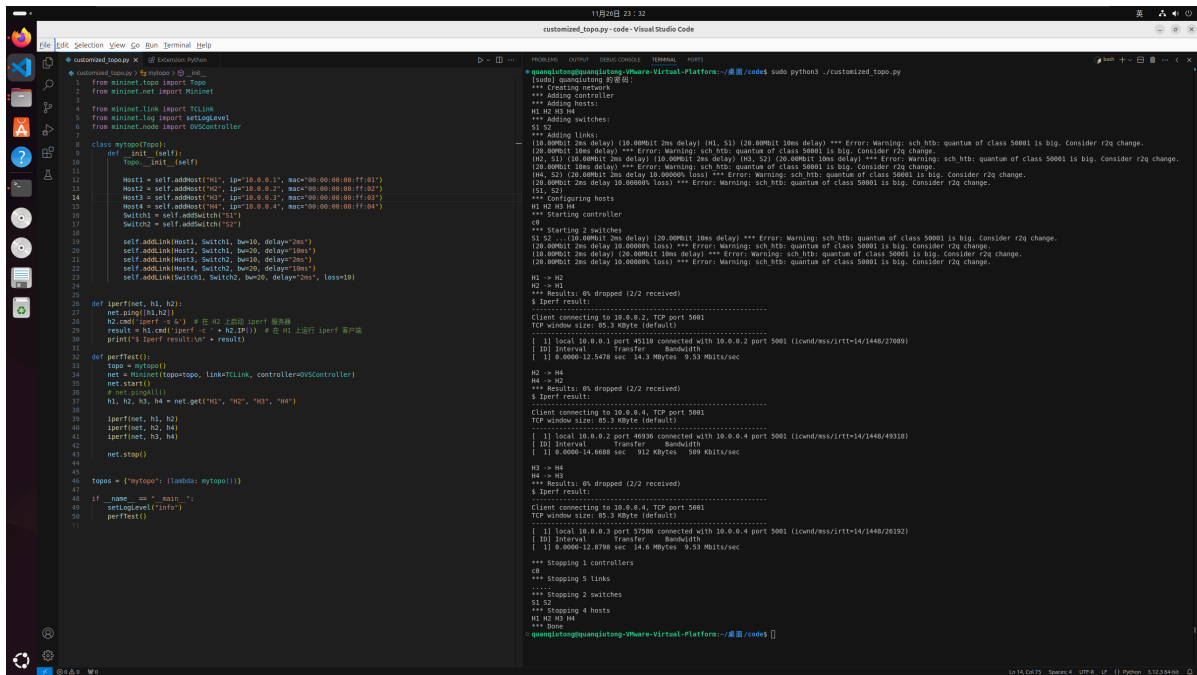
iperf version 2.1.9 (14 March 2023) pthreads

任务1：定制化拓扑

1. 利用iperf验证端到端带宽

该iperf版本下，直接用 net.iperf((h1, h2)) 会卡住，所以用了如下函数：

```
def iperf(net, h1, h2):
    net.ping([h1,h2])
    h2.cmd('iperf -s &') # 在 H2 上启动 iperf 服务器
    result = h1.cmd('iperf -c ' + h2.IP()) # 在 H1 上运行 iperf 客户端
    print("$ Iperf result:\n" + result)
```



```
customized_topo.py x Python3
1 from mininet.topo import Topo
2 from mininet.net import Mininet
3
4 from mininet.link import TCLink
5 from mininet.log import setLogLevel
6 from mininet.node import OVSController
7
8 class mytopo(Topo):
9     def __init__(self):
10         Topo.__init__(self)
11
12         Host1 = self.addHost('H1', ip='10.0.0.1', mac='00:00:00:00:00:01')
13         Host2 = self.addHost('H2', ip='10.0.0.2', mac='00:00:00:00:00:02')
14         Host3 = self.addHost('H3', ip='10.0.0.3', mac='00:00:00:00:00:03')
15         Host4 = self.addHost('H4', ip='10.0.0.4', mac='00:00:00:00:00:04')
16
17         Switch1 = self.addSwitch('S1')
18         Switch2 = self.addSwitch('S2')
19
20         self.addLink(Host1, Switch1, bw=10, delay='2ms')
21         self.addLink(Host2, Switch1, bw=10, delay='1ms')
22         self.addLink(Host3, Switch2, bw=10, delay='2ms')
23         self.addLink(Host4, Switch2, bw=10, delay='1ms')
24         self.addLink(Switch1, Switch2, bw=20, delay='2ms', loss=10)
25
26 def iperfTest():
27     net = Mininet(topo=mytopo, link=TCLink, controller=OVSController)
28     net.start()
29     # net.pingAll()
30     H1, H2, H3, H4 = net.get('H1', 'H2', 'H3', 'H4')
31
32     iperfTest(H1, H2)
33     iperfTest(H2, H3)
34     iperfTest(H3, H4)
35     net.stop()
36
37 topotest = 'mytopo'
38 if __name__ == '__main__':
39     setLogLevel('info')
40     perTest()
```

```
quantumquantum@quantum-Virtual-Platform:~/codes$ sudo python3 ./customized_topo.py
[info] quantumquant 启动成功
*** Creating network
*** Adding controller
*** Adding hosts:
H1 H2 H3 H4
*** Adding switches:
S1 S2
*** Adding links:
(10.000bit 2ms delay) (10.000bit 2ms delay) (H1, S1) (20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(H2, S1) (10.000bit 2ms delay) (10.000bit 2ms delay) (H2, S2) (20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(H3, S2) (10.000bit 2ms delay) (10.000bit 2ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(20.000bit 2ms delay) (20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
*** Configuring hosts
H1 H2 H3 H4
*** Starting controller
c0
*** Starting 2 switches
S1 S2 ... (10.000bit 2ms delay) (20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(10.000bit 2ms delay) (20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
(20.000bit 2ms delay) (20.000bit 1ms delay) *** Error: Warning: sch_htb: quantum of class 50001 is big. Consider r2q change.
*** Stopping 1 controllers
c0
*** Stopping 3 links
*** Stopping 2 switches
S1 S2
*** Stopping 4 hosts
H1 H2 H3 H4
*** Done
quantumquant@quantum-Virtual-Platform:~/codes$
```

```
$ Iperf result:
-----
Client connecting to 10.0.0.2, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 1] local 10.0.0.1 port 45110 connected with 10.0.0.2 port 5001 (icwnd/mss/irrt=14/1448/27089)
[ ID] Interval          Transfer        Bandwidth
[ 1] 0.0000-12.5478 sec  14.3 MBytes    9.53 Mb/s/sec
```

\$ Iperf result:

Client connecting to 10.0.0.2, TCP port 5001
TCP window size: 85.3 KByte (default)

[1] local 10.0.0.1 port 45110 connected with 10.0.0.2 port 5001
(icwnd/mss/irrt=14/1448/27089)
[ID] Interval Transfer Bandwidth
[1] 0.0000-12.5478 sec 14.3 MBytes 9.53 Mb/s/sec

\$ Iperf result:

Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)

[1] local 10.0.0.2 port 46936 connected with 10.0.0.4 port 5001
(icwnd/mss/irrt=14/1448/49318)
[ID] Interval Transfer Bandwidth
[1] 0.0000-14.6688 sec 912 KBytes 509 Kbits/sec

\$ Iperf result:

Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)

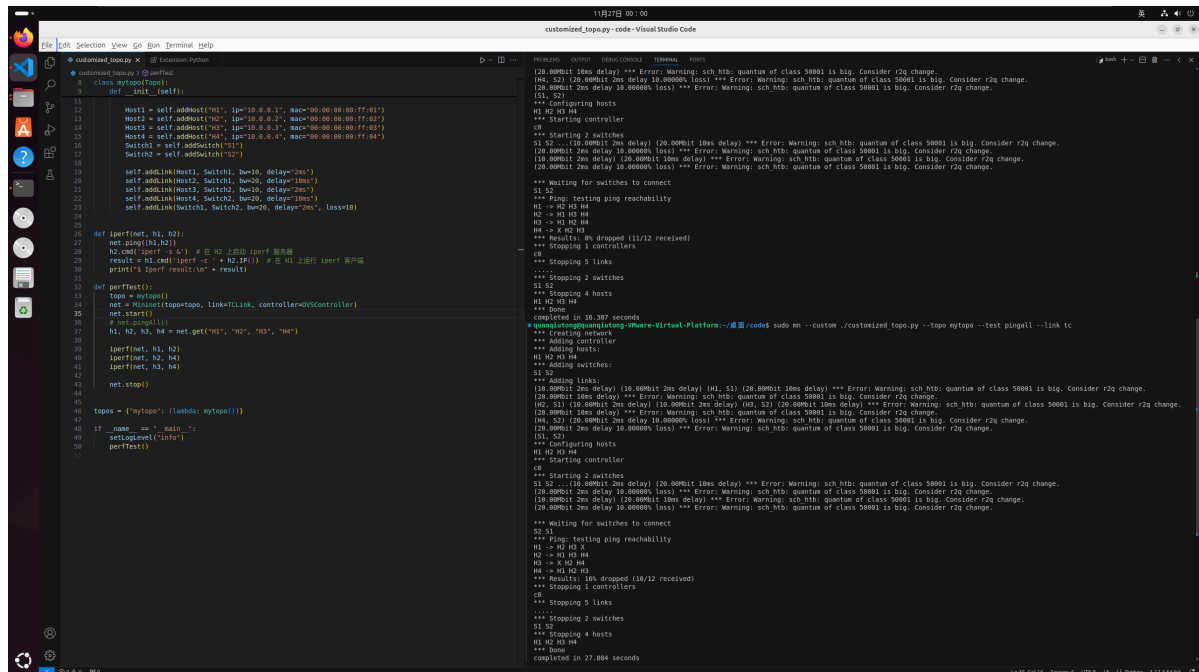
```
[ 1] local 10.0.0.3 port 57586 connected with 10.0.0.4 port 5001
(icwnd/mss/irrt=14/1448/26192)
[ ID] Interval          Transfer          Bandwidth
[ 1] 0.0000-12.8798 sec  14.6 MBytes      9.53 Mbits/sec
```

也使用要求的iperf版本（2.0.5）上的net.iperf()测试了一下，结果如下：

```
*** Iperf: testing TCP bandwidth between H1 and H2
*** Results: ['7.87 Mbits/sec', '9.13 Mbits/sec']
*** Iperf: testing TCP bandwidth between H2 and H4
*** Results: ['629 Kbits/sec', '724 Kbits/sec']
*** Iperf: testing TCP bandwidth between H3 and H4
*** Results: ['8.59 Mbits/sec', '9.93 Mbits/sec']
```

与期望的结果一致，即1->2和3->4的带宽接近理论值，2->4的带宽远低于理论值。

2. ping all 测试：



测试几次，丢包率为8%，16%，12%等，和理论值较为接近。

并且所有测试中，1->2和3->4都能ping通。

任务2：在虚拟终端上执行任务

1. Flow 1 和 Flow 2 带宽测试结果


```

[ 1] 14.5000-15.0000 sec  0.000 Bytes  0.000 bits/sec
[ 1] 15.0000-15.5000 sec  0.000 Bytes  0.000 bits/sec
[ 1] 15.5000-16.0000 sec  0.000 Bytes  0.000 bits/sec
[ 1] 16.0000-16.5000 sec  76.4 KBytes  1.25 Mbits/sec
[ 1] 16.5000-17.0000 sec  80.6 KBytes  1.32 Mbits/sec
[ 1] 17.0000-17.5000 sec  65.0 KBytes  1.07 Mbits/sec
[ 1] 17.5000-18.0000 sec  63.6 KBytes  1.04 Mbits/sec
[ 1] 18.0000-18.5000 sec   144 KBytes  2.36 Mbits/sec
[ 1] 18.5000-19.0000 sec  80.6 KBytes  1.32 Mbits/sec
[ 1] 19.0000-19.5000 sec  206 KBytes  3.37 Mbits/sec
[ 1] 19.5000-20.0000 sec  111 KBytes  1.82 Mbits/sec
[ 1] 20.0000-20.6989 sec  65.0 KBytes   762 kbits/sec
[ 1] 0.0000-20.6989 sec  3.68 MBytes  1.49 Mbits/sec

```

```

-----
Client connecting to 10.0.0.4, TCP port 5001

```

```

TCP window size: 85.3 KByte (default)
-----

```

```

[ 1] local 10.0.0.2 port 34638 connected with 10.0.0.4 port 5001
(icwnd/mss/irrt=14/1448/95447)

```

[ID]	Interval	Transfer	Bandwidth
[1]	0.0000-0.5000 sec	512 KBytes	8.39 Mbits/sec
[1]	0.5000-1.0000 sec	94.6 KBytes	1.55 Mbits/sec
[1]	1.0000-1.5000 sec	127 KBytes	2.09 Mbits/sec
[1]	1.5000-2.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	2.0000-2.5000 sec	0.000 Bytes	0.000 bits/sec
[1]	2.5000-3.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	3.0000-3.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	3.5000-4.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	4.0000-4.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	4.5000-5.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	5.0000-5.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	5.5000-6.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	6.0000-6.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	6.5000-7.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	7.0000-7.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	7.5000-8.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	8.0000-8.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	8.5000-9.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	9.0000-9.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	9.5000-10.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	10.0000-10.5000 sec	0.000 Bytes	0.000 bits/sec
[1]	10.5000-11.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	11.0000-11.5000 sec	0.000 Bytes	0.000 bits/sec
[1]	11.5000-12.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	12.0000-12.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	12.5000-13.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	13.0000-13.5000 sec	0.000 Bytes	0.000 bits/sec
[1]	13.5000-14.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	14.0000-14.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	14.5000-15.0000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	15.0000-15.5000 sec	63.6 KBytes	1.04 Mbits/sec
[1]	15.5000-16.0000 sec	0.000 Bytes	0.000 bits/sec
[1]	16.0000-16.5000 sec	0.000 Bytes	0.000 bits/sec
[1]	16.5000-17.0000 sec	63.6 KBytes	1.04 Mbits/sec

```
[ 1] 17.0000-17.5000 sec 0.000 Bytes 0.000 bits/sec
[ 1] 17.5000-18.0000 sec 0.000 Bytes 0.000 bits/sec
[ 1] 18.0000-18.5000 sec 0.000 Bytes 0.000 bits/sec
[ 1] 18.5000-19.0000 sec 0.000 Bytes 0.000 bits/sec
[ 1] 19.0000-19.5000 sec 63.6 KBytes 1.04 Mbits/sec
[ 1] 19.5000-20.0000 sec 127 KBytes 2.09 Mbits/sec
[ 1] 0.0000-26.3471 sec 2.08 MBytes 663 Kbits/sec
```

2. 观察到的带宽：

传输初期带宽较高，随后由于网络条件（丢包、拥塞控制）带宽显著下降，最终平均带宽远低于各自的理论值10Mb/s。

3. 原因

1. 链路丢包率高：交换机S1与S2之间的链路设置了10%的丢包率，导致TCP连接在传输过程中出现丢包。
2. TCP拥塞控制机制：丢包触发了TCP的拥塞控制算法，减小了拥塞窗口，降低了发送速率，导致带宽下降甚至短暂为零。
3. 带宽竞争：在10秒后，第二个TCP流开始传输，两个流共享S1-S2的20Mbps带宽，增加了网络拥塞，进一步影响了各自的吞吐量。
4. 延迟和带宽差异：主机之间的链路延迟和带宽不同，影响了TCP的性能。例如，H1到S1的延迟为2ms，而H2到S1的延迟为10ms。

4. 尝试修改 Switch S1 和 Switch S2 之间链路的丢包率，重复任务二，观察并描述在不同丢包率下出现的现象。

丢包率低时，带宽低于理论值的情况会减轻，但仍然会出现带宽下降的情况。

丢包率高时，带宽下降的情况会更加严重，甚至出现带宽为零的情况。

而且两条链路的RTT差异会导致两条链路的带宽不同，这也是导致带宽下降的原因之一。

以丢包率0%为例，结果如下：

```
-----
Client connecting to 10.0.0.3, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 1] local 10.0.0.1 port 36860 connected with 10.0.0.3 port 5001
(icwnd/mss/irrtt=14/1448/37405)
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-0.5000 sec  896 KBytes   14.7 Mbits/sec
[ 1] 0.5000-1.0000 sec  768 KBytes   12.6 Mbits/sec
[ 1] 1.0000-1.5000 sec  640 KBytes   10.5 Mbits/sec
[ 1] 1.5000-2.0000 sec  640 KBytes   10.5 Mbits/sec
[ 1] 2.0000-2.5000 sec  733 KBytes   12.0 Mbits/sec
[ 1] 2.5000-3.0000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 3.0000-3.5000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 3.5000-4.0000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 4.0000-4.5000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 4.5000-5.0000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 5.0000-5.5000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 5.5000-6.0000 sec  764 KBytes   12.5 Mbits/sec
[ 1] 6.0000-6.5000 sec  764 KBytes   12.5 Mbits/sec
```



```
[ 1] 6.5000-7.0000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 7.0000-7.5000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 7.5000-8.0000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 8.0000-8.5000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 8.5000-9.0000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 9.0000-9.5000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 9.5000-10.0000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 10.0000-10.5000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 10.5000-11.0000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 11.0000-11.5000 sec 954 KBytes 15.6 Mbits/sec
[ 1] 11.5000-12.0000 sec 1.06 MBytes 17.7 Mbits/sec
[ 1] 12.0000-12.5000 sec 637 KBytes 10.4 Mbits/sec
[ 1] 12.5000-13.0000 sec 1.12 MBytes 18.8 Mbits/sec
[ 1] 13.0000-13.5000 sec 1.18 MBytes 19.8 Mbits/sec
[ 1] 13.5000-14.0000 sec 1.18 MBytes 19.8 Mbits/sec
[ 1] 14.0000-14.5000 sec 1.37 MBytes 22.9 Mbits/sec
[ 1] 14.5000-15.0000 sec 1.37 MBytes 22.9 Mbits/sec
[ 1] 15.0000-15.5000 sec 1.24 MBytes 20.8 Mbits/sec
[ 1] 15.5000-16.0000 sec 1.43 MBytes 24.0 Mbits/sec
[ 1] 16.0000-16.5000 sec 1.49 MBytes 25.0 Mbits/sec
[ 1] 16.5000-17.0000 sec 1.55 MBytes 26.1 Mbits/sec
[ 1] 17.0000-17.5000 sec 1.62 MBytes 27.1 Mbits/sec
[ 1] 17.5000-18.0000 sec 1.62 MBytes 27.1 Mbits/sec
[ 1] 18.0000-18.5000 sec 1.24 MBytes 20.8 Mbits/sec
[ 1] 18.5000-19.0000 sec 508 KBytes 8.32 Mbits/sec
[ 1] 19.0000-19.5000 sec 509 KBytes 8.34 Mbits/sec
[ 1] 19.5000-20.0000 sec 509 KBytes 8.34 Mbits/sec
[ 1] 20.0000-33.6120 sec 127 KBytes 76.1 Kbits/sec
[ 1] 0.0000-33.6120 sec 36.9 MBytes 9.21 Mbits/sec
```

```
-----
Client connecting to 10.0.0.4, TCP port 5001
```

```
TCP window size: 85.3 KByte (default)
-----
```

```
[ 1] local 10.0.0.2 port 58586 connected with 10.0.0.4 port 5001
(icwnd/mss/irrt=14/1448/97048)
```

```
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-0.5000 sec 1.63 MBytes 27.3 Mbits/sec
[ 1] 0.5000-1.0000 sec 722 KBytes 11.8 Mbits/sec
[ 1] 1.0000-1.5000 sec 827 KBytes 13.6 Mbits/sec
[ 1] 1.5000-2.0000 sec 892 KBytes 14.6 Mbits/sec
[ 1] 2.0000-2.5000 sec 699 KBytes 11.5 Mbits/sec
[ 1] 2.5000-3.0000 sec 827 KBytes 13.6 Mbits/sec
[ 1] 3.0000-3.5000 sec 892 KBytes 14.6 Mbits/sec
[ 1] 3.5000-4.0000 sec 317 KBytes 5.20 Mbits/sec
[ 1] 4.0000-4.5000 sec 764 KBytes 12.5 Mbits/sec
[ 1] 4.5000-5.0000 sec 827 KBytes 13.6 Mbits/sec
[ 1] 5.0000-5.5000 sec 892 KBytes 14.6 Mbits/sec
[ 1] 5.5000-6.0000 sec 699 KBytes 11.5 Mbits/sec
[ 1] 6.0000-6.5000 sec 827 KBytes 13.6 Mbits/sec
[ 1] 6.5000-7.0000 sec 892 KBytes 14.6 Mbits/sec
[ 1] 7.0000-7.5000 sec 699 KBytes 11.5 Mbits/sec
[ 1] 7.5000-8.0000 sec 892 KBytes 14.6 Mbits/sec
[ 1] 8.0000-8.5000 sec 699 KBytes 11.5 Mbits/sec
```



```

[ 1] 8.5000-9.0000 sec   827 KBytes  13.6 Mbits/sec
[ 1] 9.0000-9.5000 sec   764 KBytes  12.5 Mbits/sec
[ 1] 9.5000-10.0000 sec   827 KBytes  13.6 Mbits/sec
[ 1] 10.0000-10.5000 sec   827 KBytes  13.6 Mbits/sec
[ 1] 10.5000-11.0000 sec   764 KBytes  12.5 Mbits/sec
[ 1] 11.0000-11.5000 sec   955 KBytes  15.7 Mbits/sec
[ 1] 11.5000-12.0000 sec   1.12 MBytes  18.8 Mbits/sec
[ 1] 12.0000-12.5000 sec   1.12 MBytes  18.8 Mbits/sec
[ 1] 12.5000-13.0000 sec   1.18 MBytes  19.8 Mbits/sec
[ 1] 13.0000-13.5000 sec   1.37 MBytes  22.9 Mbits/sec
[ 1] 13.5000-14.0000 sec   1.37 MBytes  22.9 Mbits/sec
[ 1] 14.0000-14.5000 sec    700 KBytes  11.5 Mbits/sec
[ 1] 14.5000-15.0000 sec   1.43 MBytes  24.0 Mbits/sec
[ 1] 15.0000-15.5000 sec    760 KBytes  12.5 Mbits/sec
[ 1] 15.5000-16.0000 sec    0.000 Bytes  0.000 bits/sec
[ 1] 16.0000-16.5000 sec   128 KBytes  2.10 Mbits/sec
[ 1] 16.5000-17.0000 sec  1020 KBytes  16.7 Mbits/sec
[ 1] 17.0000-17.5000 sec   1.18 MBytes  19.8 Mbits/sec
[ 1] 17.5000-18.0000 sec   123 KBytes  2.01 Mbits/sec
[ 1] 18.0000-18.5000 sec   638 KBytes  10.4 Mbits/sec
[ 1] 18.5000-19.0000 sec   827 KBytes  13.6 Mbits/sec
[ 1] 19.0000-19.5000 sec   1.30 MBytes  21.9 Mbits/sec
[ 1] 19.5000-20.0000 sec    508 KBytes  8.33 Mbits/sec
[ 1] 20.0000-25.8463 sec   127 KBytes   177 kbits/sec
[ 1] 0.0000-25.8463 sec  33.3 MBytes  10.8 Mbits/sec

```

- Flow 1 (H1 -> H3) 的最小带宽为 10 Mbps (受制于 H1-S1 和 S2-H3 的链路) , RTT 为 12 ms。
- Flow 2 (H2 -> H4) 的最小带宽为 20 Mbps, RTT 为 44 ms。

TCP 吞吐量大致可以用以下公式估算:

$$\text{吞吐量} = \frac{\text{TCP 窗口大小}}{\text{RTT}}$$

在 RTT 较大的情况下, TCP 的吞吐量会降低。

- 实际带宽差异原因:

Flow 1 拥有较低的 RTT (12 ms) , 因此即使链路带宽较低, TCP 也能较充分地利用带宽。

Flow 2 由于 RTT 较高 (44 ms) , 即使链路带宽更高, TCP 的吞吐量也会受到更大的限制。

TCP 协议的拥塞控制机制使得高延迟的连接增长速度较慢, 导致带宽利用率降低。

- 竞争和共享带宽:

在两个流同时进行的时间段内 (10 秒到 20 秒) , 它们需要共享 S1-S2 链路的带宽 (20 Mbps) 。

由于 Flow 1 已经建立并稳定运行, Flow 2 需要与其竞争带宽。

TCP 的公平性并不总是保证带宽平均分配, 尤其在 RTT 不同的情况下。